

The AI Policy Brief for School Leaders.

01 / The evidence

Recent trials show both promise and risk. Surveys describe use and concern. Neither validates a vendor by itself.

+0.33 g | Structured AI revision shows promise.

Four-week GCSE science trial: higher scores than self-directed revision. Preliminary: 30.7% attrition and nonstandardized tests. Harrison et al. · September 13, 2026 · Preprint

-0.37 SD | Course access alone is not enough.

University trial: lower grades in the same-course comparison. Mostly direct-answer mode, low uptake; not evidence about all tutors or K-12. Liu et al. · October 2026 · Working paper

-17% | Better practice scores can hide weaker learning.

Unrestricted AI helped practice but hurt subsequent unaided exam performance. Teacher-designed hints largely prevented the loss; they did not establish an exam gain. Bastani et al. · PNAS · June 25, 2025

84% | Faculty are concerned about original thinking.

Faculty agreed AI reduces original expression. A perception survey in higher education, not an experiment measuring learning loss. College Board · February 2026 · Figure 10

67% | Concern is rising alongside use.

Students believed schoolwork AI harms critical thinking. This measures concern, not personal cognitive decline. RAND · March 17, 2026 · Survey

62% | Most surveyed students already use it.

Homework use in December 2025, up from 48% in May. Ages 12-29 across school and college; usage is not dependence. RAND · March 17, 2026 · Survey

These studies did not evaluate Orchen. Full citations and context: orchen.ai/the-problem. Reviewed October 5, 2026.

Three honest options, and what each one costs.

Option A

Ban AI on school networks and devices.

What it buys: a clear public position and short-term board comfort. **What it costs:** students may move to personal devices. Restrictions still need clear expectations and a plan for teaching appropriate use.

Option B

Permit consumer AI with guidelines.

What it buys: honesty about reality and low cost. **What it costs:** the school is endorsing tools designed for task completion, with learning modes that vary by provider and student data on consumer terms. Guidelines without an environment to apply them in are advice, not policy. The gap between submitted work and actual understanding keeps widening, invisibly.

Option C

Provide a governed environment, and set policy around it.

What it buys: an AI the school configured to teach rather than complete work, structured visibility for faculty and parents, crisis detection that routes to a counselor, and a defensible data posture. **What it costs:** budget, an implementation effort, and understanding that a school governs its own environment rather than the open internet, which is exactly why that environment needs to be one students genuinely prefer to use. A governed tool still needs evaluation: measure independent learning and teacher workload during the pilot.

A note on who's telling you this. This brief is published by Orchen, which sells a governed AI environment, so Option C is, transparently, our business. We've tried to state all three options at their strongest. Whichever direction your school takes, the checklist below applies to any vendor you evaluate, including us.

Twelve questions to settle before any deployment.

Use this in a leadership meeting or send it to any vendor, including us. The quality of the answers tells you most of what you need to know.

01. Does the AI complete student work, or does it require student thinking before it proceeds?
Ask for a live demonstration of a student asking for the answer directly.

02. What evidence shows students can work independently after using the tool?

03. What do teachers see, and is it synthesis or surveillance?
Raw transcript access at scale is both a privacy problem and a workload problem.

04. What do parents see, and who controls that, the school or the vendor?

05. How long are student conversations retained, and is the answer specific?
"Per our privacy policy" is not a number. Get a number.

06. Is student data used to train AI models, the vendor's or their providers'?

07. What happens when a student types something concerning?
Who at the school is notified, how fast, and is there a tracked lifecycle?

08. Is staff access to student data logged, and can the school review the log?

09. How does under-13 consent actually work, and what happens without it?

10. Can the school export everything and exit cleanly?
Ask for the export format and the post-termination deletion window in writing.

11. What certifications exist today, and what does the vendor admit it doesn't have yet?
A vendor that volunteers its gaps is telling you how it will behave under pressure.

12. Can a pilot be scoped small enough that failure costs weeks, not budget?

Run the checklist against us: [Trust Center](#) and [platform documentation](#).